

GEORGIOS TRASKAS

Data Scientist | Machine Learning Engineer | Predictive Analytics

Thessaloniki, Greece | georgiost77@gmail.com | +30 6934 694457 | [linkedin.com/in/george-traskas](https://www.linkedin.com/in/george-traskas) | github.com/gtraskas

PROFESSIONAL SUMMARY

Data Scientist and ML Engineer with 8 years of experience building and deploying production machine learning systems. Specialized in predictive modeling, feature engineering, and automated data pipelines. For 4+ years at Predictive Fitness, delivered ML-driven predictive models, AWS data workflows, and production monitoring systems, reporting directly to the Head of Science and Innovation and to the Founder and CEO with weekly executive updates. Completed 90+ global consulting engagements (Upwork, 100% Job Success Score) delivering solutions across healthcare, energy, and sports science.

CORE SKILLS

Machine Learning: XGBoost, Scikit-Learn, Random Forest, Classification, Regression, Time-Series Forecasting

Deep Learning: Custom model training, TensorFlow, PyTorch, Embedded deployment

Languages & Tools: Python, SQL, R, Pandas, NumPy, Jupyter, Matplotlib, Seaborn, Plotly

Cloud & Pipelines: AWS (Lambda, ECS Fargate, Athena, S3), Docker, GitHub Actions CI/CD, ETL design

MLOps & Data Architecture: Data warehousing, data modeling, pipeline orchestration, model monitoring, drift detection

Databases: PostgreSQL, pgvector, ChromaDB, Supabase, DynamoDB, AWS Athena

AI & Agents: RAG, LLM Fine-Tuning, MCP (Model Context Protocol), Multi-Agent Orchestration, Ensemble Modeling

Monitoring & Visualization: Grafana (custom dashboards, alerting), Sentry (tracing, profiling), production reliability tracking

Analytics & Reporting: Statistical modeling, hypothesis testing, cohort analysis, causal inference (do-calculus, Pearl)

PROFESSIONAL EXPERIENCE

Predictive Fitness, Inc. | Data Scientist & ML Engineer | May 2022 – Present (part-time since May 2026) | Remote

- Reported directly to the Head of Science and Innovation and to the Founder and CEO, presenting model performance updates and data-driven recommendations in weekly executive meetings.
- Optimized multi-source data ingestion pipelines from wearables and fitness devices on AWS (Lambda, ECS Fargate, Athena, GraphQL APIs), developed science-based physiological models using exercise physiology equations and empirical research, and researched and built production ML models for workout outcome prediction and performance forecasting.
- Applied causal inference (Pearl's framework, DAGs, do-calculus) to trial conversion prediction helping the marketing team to directly action targeting campaigns and user nudges.
- Integrated GraphQL APIs with Python backend workflows for real-time feature engineering and data ingestion.
- Built the company's MCP (Model Context Protocol) server exposing athlete data as tools for AI clients: Python proof-of-concept productionized in TypeScript with Prisma over GraphQL and a RAG layer for training-metric knowledge.
- Redesigned production alerting architecture using Grafana and Sentry, reducing false alert noise by over 90% from dozens of daily false positives to fewer than five genuine alerts per week.

Data Science Solutions | Independent AI & Data Science Consultant | Jul 2020 – Present | Remote

- Completed 90+ client engagements globally via Upwork with a 100% Job Success Score, delivering predictive models, ML pipelines, and data analytics solutions across healthcare, energy, and retail sectors.
- MedTech (Israel):** Trained and optimized a deep learning classification model for embedded microcontroller deployment, achieving 97% accuracy with TFLite quantization for hardware-constrained production environments.
- French Institute of Gynaecology:** Predicted IUI fertility outcomes on an imbalanced clinical dataset (8.5% positive class) using SMOTE and repeated stratified K-fold, evaluated on recall and F1 instead of misleading accuracy; deployed as a clinical decision-support tool.

Centre for Research & Technology Hellas | Chemist & Automation Engineer | Dec 2007 – Jul 2022 | Greece

- Automated analytical data pipelines, increasing system throughput by 80% and reducing delivery times by 50%.
- Applied statistical analysis and experimental design to research projects, delivering reporting tools consumed by scientific teams and management.

SELECTED PROJECTS

MolekitChen – AI/RAG Food-Science iOS App | 2026, live on **App Store**

- Designed and deployed a full-stack AI mobile application parsing peer-reviewed food-science papers with live citations. Built on FastAPI, Gemini AI, Supabase (PostgreSQL + pgvector), and Swift (iOS). Deployed on Render, live on the App Store.

Wine-VFM – Autonomous Multi-Agent Bargain Hunter | 2026, live demo on **Modal**

- Built a multi-agent system estimating wine value-for-money from tasting notes with three models – a fine-tuned Llama 3.2 3B, a RAG estimator over ~88K embedded notes, and a neural network – blended with regression-fitted ensemble weights (evaluated with MAE/R²) under an autonomous LLM planner that scans a live wine shop and sends push alerts.
- **Stack:** Python, Modal (serverless GPU), PyTorch, ChromaDB, Scikit-Learn, Gradio. Code: github.com/gtraskas/wine-vfm

EDUCATION

MSc, Chemical Technology | Aristotle University of Thessaloniki | 2004

BSc, Chemistry | Aristotle University of Thessaloniki | 2002

CERTIFICATIONS

IBM Generative AI (2026) • Claude AI, Anthropic (2026) • AutoML (2022) • AWS Cloud Essentials (2022)

ML Engineer Nanodegree, Udacity (2020) • Google Big Data & ML Fundamentals (2019) • TensorFlow Developer (2019) • Stanford ML, Coursera (2018) • Data Analyst, Udacity (2017)